

hw6_MachineLearning_fall2018

October 7, 2018

1 Data-X Fall 2018: Homework 06

1.0.1 Machine Learning

Authors: Sana Iqbal (Part 1, 2, 3)

In this homework, you will do some exercises with prediction.

```
In [1]: import numpy as np
import pandas as pd

In [2]: # machine learning libraries
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
import xgboost as xgb
```

1.1 Part 1

__ 1. Read `diabetesdata.csv` file into a pandas dataframe. About the data: __

1. **TimesPregnant:** Number of times pregnant
2. **glucoseLevel:** Plasma glucose concentration a 2 hours in an oral glucose tolerance test
3. **BP:** Diastolic blood pressure (mm Hg)
4. **insulin:** 2-Hour serum insulin (mu U/ml)
5. **BMI:** Body mass index (weight in kg/(height in m)²)
6. **pedigree:** Diabetes pedigree function
7. **Age:** Age (years)
8. **IsDiabetic:** 0 if not diabetic or 1 if diabetic)

```
In [3]: #Read data & print it
data = pd.read_csv("diabetesdata.csv")
```

```
data.info()
data.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 8 columns):
TimesPregnant      768 non-null int64
glucoseLevel       734 non-null float64
BP                 768 non-null int64
insulin            768 non-null int64
BMI                768 non-null float64
Pedigree           768 non-null float64
Age                735 non-null float64
IsDiabetic         768 non-null int64
dtypes: float64(4), int64(4)
memory usage: 48.1 KB
```

```
Out[3]:
```

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic
0	6	148.0	72	0	33.6	0.627	50.0	1
1	1	NaN	66	0	26.6	0.351	31.0	0
2	8	183.0	64	0	23.3	0.672	NaN	1
3	1	NaN	66	94	28.1	0.167	21.0	0
4	0	137.0	40	168	43.1	2.288	33.0	1

2. Calculate the percentage of NaN values in each column.

```
In [4]: NullsPerColumn = pd.DataFrame()
NullsPerColumn['Percentage Null'] = data.isna().sum() / data.shape[0]
NullsPerColumn
```

```
Out[4]:
```

	Percentage Null
TimesPregnant	0.000000
glucoseLevel	0.044271
BP	0.000000
insulin	0.000000
BMI	0.000000
Pedigree	0.000000
Age	0.042969
IsDiabetic	0.000000

```
In [5]: ###RUN THIS CELL BUT DO NOT ALTER IT
assert all(NullsPerColumn.columns == ['Percentage Null'])
assert NullsPerColumn['Percentage Null'][-2] == 0.04296875
```

3. Calculate the TOTAL percent of ROWS with NaN values in the dataframe (make sure values are floats).

```
In [6]: PercentNull = data.isna().max(axis="columns").sum()/data.shape[0]
PercentNull
```

```
Out[6]: 0.08333333333333333
```

4. Split data into train_df and test_df with 15% test split.

```
In [7]: #split values
        from sklearn.model_selection import train_test_split
        train_df, test_df = train_test_split(data, test_size=0.15)
```

```
In [8]: ###RUN THIS CELL BUT DO NOT ALTER IT
        np.testing.assert_almost_equal(float(len(train_df))/float(len(data)), 0.8489583333333333)
        np.testing.assert_almost_equal(float(len(test_df))/float(len(data)), 0.15104166666666666)
```

5. Replace the Nan values in train_df and test_df with the mean of EACH feature.

```
In [9]: train_df = train_df.fillna(train_df.mean())
        test_df = test_df.fillna(test_df.mean())
```

```
In [10]: ###RUN THIS CELL BUT DO NOT ALTER IT
         assert sum(train_df.isnull().sum()) == 0
         assert sum(test_df.isnull().sum()) == 0
```

6. Split train_df & test_df into X_train, Y_train and X_test, Y_test. Y_train and Y_test should only have the column we are trying to predict, IsDiabetic.

```
In [11]: X_train = train_df.drop(columns="IsDiabetic")
        Y_train = train_df["IsDiabetic"]
        X_test = test_df.drop(columns="IsDiabetic")
        Y_test = test_df["IsDiabetic"]
```

```
In [12]: ###RUN THIS CELL BUT DO NOT ALTER IT
         assert [X_train.shape, Y_train.shape, X_test.shape, Y_test.shape] == [(652, 7), (652,)]
```

7. Use this dataset to train perceptron, logistic regression and random forest models using 15% test split. Report training and test accuracies.

```
In [13]: # Logistic Regression

        logreg = LogisticRegression()
        logreg.fit(X_train, Y_train)
        logreg_train_acc = logreg.score(X_train, Y_train)
        logreg_test_acc = logreg.score(X_test, Y_test)
        print('logreg training accuracy= ', logreg_train_acc)
        print('logreg test accuracy= ', logreg_test_acc)
```

```
logreg training accuracy= 0.7745398773006135
logreg test accuracy= 0.7327586206896551
```

```
c:\users\resident\documents\data\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning
FutureWarning)
```

```
In [14]: # Perceptron
```

```
perceptron = Perceptron()
perceptron.fit(X_train, Y_train)
perceptron_train_acc = perceptron.score(X_train, Y_train)
perceptron_test_acc = perceptron.score(X_test, Y_test)
print('perceptron training accuracy= ',perceptron_train_acc)
print('perceptron test accuracy= ',perceptron_test_acc)
```

```
perceptron training accuracy= 0.5521472392638037
perceptron test accuracy= 0.5689655172413793
```

```
c:\users\resident\documents\data\lib\site-packages\sklearn\linear_model\stochastic_gradient.py
FutureWarning)
```

```
In [15]: # Adaboost
```

```
adaboost = AdaBoostClassifier()
adaboost.fit(X_train, Y_train)
adaboost_train_acc = adaboost.score(X_train, Y_train)
adaboost_test_acc = adaboost.score(X_test, Y_test)
print('adaboost training accuracy= ',adaboost_train_acc)
print('adaboost test accuracy= ',adaboost_test_acc)
```

```
adaboost training accuracy= 0.8220858895705522
adaboost test accuracy= 0.6810344827586207
```

```
In [16]: # Random Forest
```

```
random_forest = RandomForestClassifier()
random_forest.fit(X_train, Y_train)
random_forest_train_acc = random_forest.score(X_train, Y_train)
random_forest_test_acc = random_forest.score(X_test, Y_test)
print('random_forest training accuracy= ',random_forest_train_acc)
print('random_forest test accuracy= ',random_forest_test_acc)
```

```
random_forest training accuracy= 0.9938650306748467
random_forest test accuracy= 0.7586206896551724
```

```
c:\users\resident\documents\data\lib\site-packages\sklearn\ensemble\forest.py:248: FutureWarning
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```

8. Is mean imputation is the best type of imputation to use? Why or why not? What are some other ways to impute the data?

Mean imputation is one of the easiest imputation to use. It is not the best type of imputation since it creates a bias for this value of the feature. There are many other ways to do imputation,

among others: - mean imputation from a specific number of groups (create N groups so that each row belongs to one group, for each row, replace nan values by the mean of the feature in all the examples of the group) - create an estimator (like linear regression) to predict the missing features with the others.

1.2 Part 2

1. Add columns BMI_band__ & Pedigree_band to Data by cutting BMI & Pedigree into 3 intervals. PRINT the first 5 rows of __data.

```
In [17]: data['BMI_band'] = pd.cut(data['BMI'], 3)
        data['Pedigree_band'] = pd.cut(data['Pedigree'], 3)
        data.head()
```

```
Out[17]:
```

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic	\
0	6	148.0	72	0	33.6	0.627	50.0	1	
1	1	NaN	66	0	26.6	0.351	31.0	0	
2	8	183.0	64	0	23.3	0.672	NaN	1	
3	1	NaN	66	94	28.1	0.167	21.0	0	
4	0	137.0	40	168	43.1	2.288	33.0	1	

	BMI_band	Pedigree_band
0	(22.367, 44.733]	(0.0757, 0.859]
1	(22.367, 44.733]	(0.0757, 0.859]
2	(22.367, 44.733]	(0.0757, 0.859]
3	(22.367, 44.733]	(0.0757, 0.859]
4	(22.367, 44.733]	(1.639, 2.42]

1a. Print the category intervals for BMI_band__ & __Pedigree_band.

```
In [18]: print('BMI_Band_Interval: {}'.format(data['BMI_band'].unique()))
        print('Pedigree_Band_Interval: {}'.format(data['Pedigree_band'].unique()))
```

BMI_Band_Interval: [(22.367, 44.733], (-0.0671, 22.367], (44.733, 67.1]]
Categories (3, interval[float64]): [(-0.0671, 22.367] < (22.367, 44.733] < (44.733, 67.1]]
Pedigree_Band_Interval: [(0.0757, 0.859], (1.639, 2.42], (0.859, 1.639]]
Categories (3, interval[float64]): [(0.0757, 0.859] < (0.859, 1.639] < (1.639, 2.42]]

2. Group data__ by Pedigree_band & determine ratio of diabetic in each band.__

```
In [19]: pedigree_DiabeticRatio = data.groupby("Pedigree_band", as_index=False).mean()
        pedigree_DiabeticRatio["IsDiabetic"]
```

```
Out[19]:
```

Pedigree_band	IsDiabetic
0	0.327007
1	0.540541
2	0.444444

Name: IsDiabetic, dtype: float64

2a. Group data__ by BMI_band & determine ratio of diabetic in each band.__

```
In [20]: BMI_DiabeticRatio = data.groupby("BMI_band", as_index=False).mean()
        BMI_DiabeticRatio["IsDiabetic"]
```

```
Out[20]: 0    0.039216
         1    0.358297
         2    0.611111
        Name: IsDiabetic, dtype: float64
```

```
In [21]: ###RUN THIS CELL BUT DO NOT ALTER IT
        assert BMI_DiabeticRatio['IsDiabetic'][1] == 0.35829662261380324
        assert pedigree_DiabeticRatio['IsDiabetic'][1] == 0.5405405405405406
```

3. Convert these features - 'BP','insulin','BMI' and 'Pedigree' into categorical values by mapping different bands of values of these features to integers 0,1,2.

HINT: USE pd.cut with bin=3 to create 3 bins

```
In [22]: list_features = ['BP', 'insulin', 'BMI', 'Pedigree']
        for feature in list_features:
            bins = pd.cut(data[feature], 3, labels=False)
            data[feature] = bins

        data = data.drop(columns=["BMI_band", "Pedigree_band"])
```

```
In [23]: ###RUN THIS CELL BUT DO NOT ALTER IT
        assert sum(data['insulin'])==49
        assert sum(data['BMI'])==753
        assert sum(data['Pedigree'])==92
```

4. Now consider the original dataset again, instead of generalizing the NAN values with the mean of the feature we will try assigning values to NANs based on some hypothesis. For example for age we assume that the relation between BMI and BP of people is a reflection of the age group. We can have 9 types of BMI and BP relations and our aim is to find the median age of each of that group:

Your Age guess matrix will look like this:

BMI	0	1	2
BP			
0	a00	a01	a02
1	a10	a11	a12
2	a20	a21	a22

Create a guess_matrix for NaN values of 'Age' (using 'BMI' and 'BP') and 'glucoseLevel' (using 'BP' and 'Pedigree') for the given dataset and assign values accordingly to the NaNs in 'Age' or 'glucoseLevel'.

Refer to how we guessed age in the titanic notebook in the class.

```
In [24]: age_matrix, glucose_matrix = np.zeros((3, 3)), np.zeros((3, 3))
        age_df, glucose_df = pd.DataFrame(), pd.DataFrame()
```

```

matrices = [age_matrix, glucose_matrix]
dataframes = [age_df, glucose_df]
columns = ["Age", "glucoseLevel"]
group_features = ["BMI", "Pedigree"]
for matrix, df, column, feature in zip(matrices, dataframes, columns, group_features):
    for i in range(0, 3):
        for j in range(0, 3):
            guess_df = data[(data['BP'] == i) \
                            &(data[feature] == j)][column].dropna()

            # Extract the median age for this group
            # (less sensitive) to outliers
            guess_value = guess_df.median()

            # Convert random age float to int
            matrix[i, j] = int(guess_value)

df = pd.DataFrame(matrix)
df.columns.name = feature
df.index.name = 'BP'
print("-"*35)
print('Guess table for {}: \n'.format(column), df)

print (' \nAssigning median {value} values to NAN {value} values in the dataset...')
print()
for i in range(0, 3):
    for j in range(0, 3):
        data.loc[ (data[column].isnull()) & (data['BP'] == i) \
                  & (data[feature] == j), column] = matrix[i,j]

data[column] = data[column].astype(int)

```

Guess table for Age:

BMI	0	1	2
BP			
0	24.0	29.0	33.0
1	25.0	29.0	32.0
2	55.0	37.0	31.0

Assigning median Age values to NAN Age values in the dataset...

Guess table for glucoseLevel:

Pedigree	0	1	2
BP			

0	115.0	127.0	137.0
1	112.0	115.0	149.0
2	133.0	129.0	159.0

Assigning median glucoseLevel values to NAN glucoseLevel values in the dataset...

5. Now, convert 'glucoseLevel' and 'Age' features also to categorical variables of 4 categories each. PRINT the head of data__

```
In [25]: list_features = ['glucoseLevel', 'Age']
        for feature in list_features:
            bins = pd.cut(data[feature], 4, labels=False)
            data[feature] = bins
        data.head()
```

```
Out[25]:
```

	TimesPregnant	glucoseLevel	BP	insulin	BMI	Pedigree	Age	IsDiabetic
0	6	2	1	0	1	0	1	1
1	1	2	1	0	1	0	0	0
2	8	3	1	0	1	0	0	1
3	1	2	1	0	1	0	0	0
4	0	2	0	0	1	2	0	1

6. Use this dataset (with all features in categorical form) to train perceptron, logistic regression and random forest models using 15% test split. Report training and test accuracies.

```
In [26]: train_df, test_df = train_test_split(data, test_size=0.15)
        X_train = train_df.drop(columns="IsDiabetic")
        Y_train = train_df["IsDiabetic"]
        X_test = test_df.drop(columns="IsDiabetic")
        Y_test = test_df["IsDiabetic"]
        X_train.shape, Y_train.shape, X_test.shape
```

```
Out[26]: ((652, 7), (652,), (116, 7))
```

```
In [27]: # Logistic Regression
        logreg = LogisticRegression()
        logreg.fit(X_train, Y_train)
        logreg_train_acc = logreg.score(X_train, Y_train)
        logreg_test_acc = logreg.score(X_test, Y_test)
        print('logreg training accuracy= ', logreg_train_acc)
        print('logreg test accuracy= ', logreg_test_acc)
```

```
logreg training accuracy= 0.7515337423312883
logreg test accuracy= 0.6896551724137931
```

```
c:\users\resident\documents\data\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning
```



```
In [28]: # Perceptron
```

```
perceptron = Perceptron()
perceptron.fit(X_train, Y_train)
perceptron_train_acc = perceptron.score(X_train, Y_train)
perceptron_test_acc = perceptron.score(X_test, Y_test)
print('perceptron training accuracy= ',perceptron_train_acc)
print('perceptron test accuracy= ',perceptron_test_acc)
```

```
perceptron training accuracy= 0.7024539877300614
```

```
perceptron test accuracy= 0.6896551724137931
```

```
c:\users\resident\documents\data\lib\site-packages\sklearn\linear_model\stochastic_gradient.py
FutureWarning)
```

```
In [95]: # Random Forest
```

```
random_forest = RandomForestClassifier()
random_forest.fit(X_train, Y_train)
random_forest_train_acc = random_forest.score(X_train, Y_train)
random_forest_test_acc = random_forest.score(X_test, Y_test)
print('random_forest training accuracy= ',random_forest_train_acc)
print('random_forest test accuracy= ',random_forest_test_acc)
```

```
random_forest training accuracy= 0.8696319018404908
```

```
random_forest test accuracy= 0.7155172413793104
```

```
c:\users\resident\documents\data\lib\site-packages\sklearn\ensemble\forest.py:248: FutureWarning
"10 in version 0.20 to 100 in 0.22.", FutureWarning)
```